

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras import layers
from tensorflow.keras.utils import get_file
import time
import os
from scipy.io import loadmat
%load_ext tensorboard
```

## Loading SVHN Data

```
In [2]: data_file_train = get_file('train_32x32.mat', origin='http://ufldl.stanford.edu/housenumbers/train_32x32.mat')
data_file_valid = get_file('test_32x32.mat', origin='http://ufldl.stanford.edu/housenumbers/test_32x32.mat')

train_dir = os.path.join(os.path.dirname(data_file_train), 'train_32x32.mat')
valid_dir = os.path.join(os.path.dirname(data_file_valid), 'test_32x32.mat')

train_data = loadmat(train_dir)
valid_data = loadmat(valid_dir)

X_train, y_train = train_data['X'], train_data['Y']
X_valid, y_valid = valid_data['X'], valid_data['Y']

print(X_train.shape) # shape is in (image_height, image_width, channels, records)
```

```
Downloading data from http://ufldl.stanford.edu/housenumbers/train_32x32.mat
182042624/182040794 [=====] - 2s 0us/step
Downloading data from http://ufldl.stanford.edu/housenumbers/test_32x32.mat
64282624/64275384 [=====] - 1s 0us/step
(32, 32, 3, 73257)
```

```
In [3]: X_train = np.rollaxis(X_train, 3)
X_valid = np.rollaxis(X_valid, 3)

print(X_train.shape) # convert shape to (records, image_height, image_width, channels)
print(X_valid.shape)

(73257, 32, 32, 3)
(26032, 32, 32, 3)
```

```
In [4]: # glimpse of the real images
sample_images = X_train[np.random.choice(len(X_train), size=64, replace=False)]

plt.figure(figsize=(8, 8))
for i in range(64):
    plt.subplot(8, 8, i+1)
    plt.imshow(sample_images[i])
    plt.xticks([])
    plt.yticks([])
plt.tight_layout()
plt.show()
```



## Preprocessing the data

```
In [5]: def preprocess(x):
    return (x-127.5)/127.5 # standardize to [-1, 1], so that tanh function is applicable

def decode(x):
    return np.uint8(127.5*x+127.5) # make sure to use uint8 type otherwise the image won't display properly
```

```
In [6]: X_train_real = preprocess(X_train)
X_valid_real = preprocess(X_valid)
```

## Generator

```
In [7]: def build_generator(input_size, leaky_alpha=0.2):
    '''generates images in (32,32,3)
    with up-sampling technique'''

    model = tf.keras.Sequential([
        layers.Dense(4*4*512, input_shape=(input_size,)),
        layers.Reshape(target_shape=(4, 4, 512)), # 4,4,512
        layers.BatchNormalization(),
        layers.LeakyReLU(alpha=leaky_alpha),
        layers.Dropout(0.2),
        layers.Conv2DTranspose(256, kernel_size=5, strides=2, padding='same', use_bias=False), # 8,8,256
        layers.BatchNormalization(),
        layers.LeakyReLU(alpha=leaky_alpha),
        layers.Conv2DTranspose(128, kernel_size=5, strides=2, padding='same', use_bias=False), # 6,16,128
        layers.BatchNormalization(),
        layers.LeakyReLU(alpha=leaky_alpha),
        layers.Conv2DTranspose(3, kernel_size=5, strides=2, padding='same', use_bias=False), # 2,32,3
        layers.Activation('tanh')
    ])

    return model
```

```
In [8]: g = build_generator(100, 0.2)
g.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
dense (Dense)	(None, 8192)	827392
reshape (Reshape)	(None, 4, 4, 512)	0
batch_normalization (BatchNo	(None, 4, 4, 512)	2048
leaky_re_lu (LeakyReLU)	(None, 4, 4, 512)	0
dropout (Dropout)	(None, 4, 4, 512)	0
conv2d_transpose (Conv2DTran	(None, 8, 8, 256)	3276800
batch_normalization_1 (Batch	(None, 8, 8, 256)	1024
leaky_re_lu_1 (LeakyReLU)	(None, 8, 8, 256)	0
conv2d_transpose_1 (Conv2DTr	(None, 16, 16, 128)	819200
batch_normalization_2 (Batch	(None, 16, 16, 128)	512
leaky_re_lu_2 (LeakyReLU)	(None, 16, 16, 128)	0
conv2d_transpose_2 (Conv2DTr	(None, 32, 32, 3)	9600
activation (Activation)	(None, 32, 32, 3)	0
<hr/>		
Total params:	4,936,576	
Trainable params:	4,934,784	
Non-trainable params:	1,792	

## Discriminator

A classifier to distinguish between real and fake images.

Input: 32x32x3 (32x32 image with 3 channels), with values between [-1,1]

Output: probability of image being real

```
In [9]: def build_discriminator(leaky_alpha=0.2):
    model = tf.keras.Sequential([
        layers.Conv2D(64, (5,5), strides=2, padding='same', input_shape=(32,32,3)), # 16,16,64
        layers.LeakyReLU(leaky_alpha),
        layers.Dropout(0.2),
        layers.Conv2D(128, (5,5), strides=2, padding='same'), # 8,8,128
        layers.BatchNormalization(),
        layers.LeakyReLU(leaky_alpha),
        layers.Dropout(0.2),
        layers.Conv2D(256, (5,5), strides=2, padding='same'), # 4,4,256
        layers.BatchNormalization(),
        layers.LeakyReLU(leaky_alpha),
        layers.Flatten(),
        layers.Dense(1, activation='sigmoid')
    ])
    return model
```

```
In [10]: d = build_discriminator(0.2)
d.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 16, 16, 64)	4864
leaky_re_lu_3 (LeakyReLU)	(None, 16, 16, 64)	0
dropout_1 (Dropout)	(None, 16, 16, 64)	0
conv2d_1 (Conv2D)	(None, 8, 8, 128)	204928
batch_normalization_3 (Batch Normalization)	(None, 8, 8, 128)	512
leaky_re_lu_4 (LeakyReLU)	(None, 8, 8, 128)	0
dropout_2 (Dropout)	(None, 8, 8, 128)	0
conv2d_2 (Conv2D)	(None, 4, 4, 256)	819456
batch_normalization_4 (Batch Normalization)	(None, 4, 4, 256)	1024
leaky_re_lu_5 (LeakyReLU)	(None, 4, 4, 256)	0
flatten (Flatten)	(None, 4096)	0
dense_1 (Dense)	(None, 1)	4097

Total params: 1,034,881  
Trainable params: 1,034,113  
Non-trainable params: 768

## DCGAN

```
In [11]: def build_gan(input_size, leaky_alpha,
                    discriminator_lr, generator_lr,
                    dis_beta_1, gen_beta_1):

    ''' input_size: the length of vector that would be put in the generator
        beta_1: the exponential decay rate for the 1st moment estimates in Adam optimizer
    '''

    # generator
    generator = build_generator(input_size, leaky_alpha)

    # discriminator
    discriminator = build_discriminator(leaky_alpha)
    discriminator.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=discriminator_lr, beta_1=dis_beta_1),
                           loss='binary_crossentropy')

    # GAN
    gan = tf.keras.Sequential([generator, discriminator])
    gan.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=generator_lr, beta_1=gen_beta_1),
                loss='binary_crossentropy')

    return gan, generator, discriminator
```

## Training GAN

```
In [12]: # generate random vectors to be put into generator
def generate_random_vectors(n_samples, input_size):
    return np.random.normal(loc=0, scale=1, size=(n_samples, input_size))

# set discriminator to trainable or not
def set_to_trainable(model, trainable):
    for layer in model.layers:
        layer.trainable = trainable

# show 64 generated images in square
def show_images(generated_images):
    n_images = len(generated_images)
    cols = 8
    rows = n_images//cols

    plt.figure(figsize=(8, 8))
    for i in range(n_images):
        img = decode(generated_images[i])
        ax = plt.subplot(rows, cols, i+1)
        plt.imshow(img)
        plt.xticks([])
        plt.yticks([])
    plt.tight_layout()
    plt.show()
```

To train a GAN:

- train discriminator with real images with the label smoothing (labels=1-smooth) (**smooth: trick**)
- train discriminator with fake images generated by the generator (labels=0)
- set the discriminator non-trainable
- train the GAN with generated images (labels=1) (i.e. train the generator with discriminator fixed)

```
In [13]: # Transform train_on_batch return value to dict
def named_logs(model, logs):
    result = {}
    for l in zip(['training_loss', 'validation_loss'], logs):
        result[l[0]] = l[1]
    return result
```

```
In [14]: # the seed would be repeatedly used for image output at different epoch
np.random.seed(2020)
seed = np.random.normal(loc=0, scale=1, size=(64, 100))
```

```
In [15]: !rm -rf ./logs_SVHN/
```

```
In [16]: def train(generator_lr,
             discriminator_lr,
             gen_beta_1,
             dis_beta_1,
             leaky_alpha,
             smooth=0.1,
             input_size=100,
             epochs=50,
             batch_size=128,
             val_size=36):

    start = time.time()

    # labels for the batch size and the validation size
    y_train_real, y_train_fake = np.ones([batch_size, 1]), np.zeros([batch_size, 1])
    y_val_real, y_val_fake = np.ones([val_size, 1]), np.zeros([val_size, 1])

    # create a GAN, a generator and a discriminator
    gan, generator, discriminator = build_gan(input_size, leaky_alpha, discriminator_lr, generator_lr, dis_beta_1, gen_beta_1)

    # prepare for tensorboard
    tb_dis = tf.keras.callbacks.TensorBoard(
        log_dir='./logs_SVHN/discriminator',
        histogram_freq=1,
        write_graph=True)
    tb_dis.set_model(discriminator)

    tb_gan = tf.keras.callbacks.TensorBoard(
        log_dir='./logs_SVHN/gan',
        histogram_freq=1,
        write_graph=True)
    tb_gan.set_model(gan)

    # iterate epochs to train the GAN
    num_batches = len(X_train_real)//batch_size
    for e in range(epochs):
        for i in range(num_batches):
            # real images in a batch
            X_batch_real = X_train_real[i*batch_size:(i+1)*batch_size]

            # generate a batch of fake images using random vectors
            random_vectors = generate_random_vectors(batch_size, input_size)
            X_batch_fake = generator.predict_on_batch(random_vectors)

            # train the discriminator
            # try: soft labels - random number between 0 and 0.1 to represent 0 labels (real images)
            #       random number between 0.9 and 1.0 to represent 1 labels
            #       smooth = np.random.uniform(low=0,high=0.1,size=1)[0]
            set_to_trainable(discriminator, True)
            dis_loss_train = discriminator.train_on_batch(X_batch_real, y_train_real-0.1)
            dis_loss_train += discriminator.train_on_batch(X_batch_fake, y_train_fake)

            # train the GAN (i.e. train the generator while keeping discriminator fixed)
            set_to_trainable(discriminator, False)
            # set the target of fake images to 1 because we want the loss between fake images and 1 (label for real image) to be minimized
            # i.e. we want the fake images to be real
            gen_loss_train = gan.train_on_batch(random_vectors, y_train_real)

            # keep noticed by the training process
            if (i+1) % 50 == 0:
                print('This is the {}/{} of epoch {}'.format(i+1, num_batches, e+1))

            # compute loss on validation set
            X_val_real = X_valid_real[np.random.choice(len(X_valid_real)), size=val_size, replace=False]

            random_vectors = generate_random_vectors(val_size, input_size)
            X_val_fake = generator.predict_on_batch(random_vectors)

            dis_loss_val = discriminator.test_on_batch(X_val_real, y_val_real)
            dis_loss_val += discriminator.test_on_batch(X_val_fake, y_val_fake)
            gen_loss_val = gan.test_on_batch(random_vectors, y_val_real)

            print("Epoch: {}/{} Discriminator Loss: {:.4f} Generator Loss: {:.4f}".format(e+1, epochs,
dis_loss_val, gen_loss_val))

            tb_dis.on_epoch_end(e, named_logs(discriminator, [dis_loss_train, dis_loss_val]))
            tb_gan.on_epoch_end(e, named_logs(gan, [gen_loss_train, gen_loss_val]))

        if e==0 or (e+1) % 5 == 0:
```

```
    imgs = generator.predict_on_batch(seed)
    show_images(imgs)

end = time.time()
print(end-start)
return generator
```

```
In [17]: train(generator_lr=0.0001,
           discriminator_lr=0.001,
           gen_beta_1=0.5,
           dis_beta_1=0.5,
           leaky_alpha=0.2,
           smooth=0.1,
           input_size=100,
           epochs=200,
           batch_size=128,
           val_size=36)
```

This is the 50/572 of epoch 1.

This is the 100/572 of epoch 1.

This is the 150/572 of epoch 1.

This is the 200/572 of epoch 1.

This is the 250/572 of epoch 1.

This is the 300/572 of epoch 1.

This is the 350/572 of epoch 1.

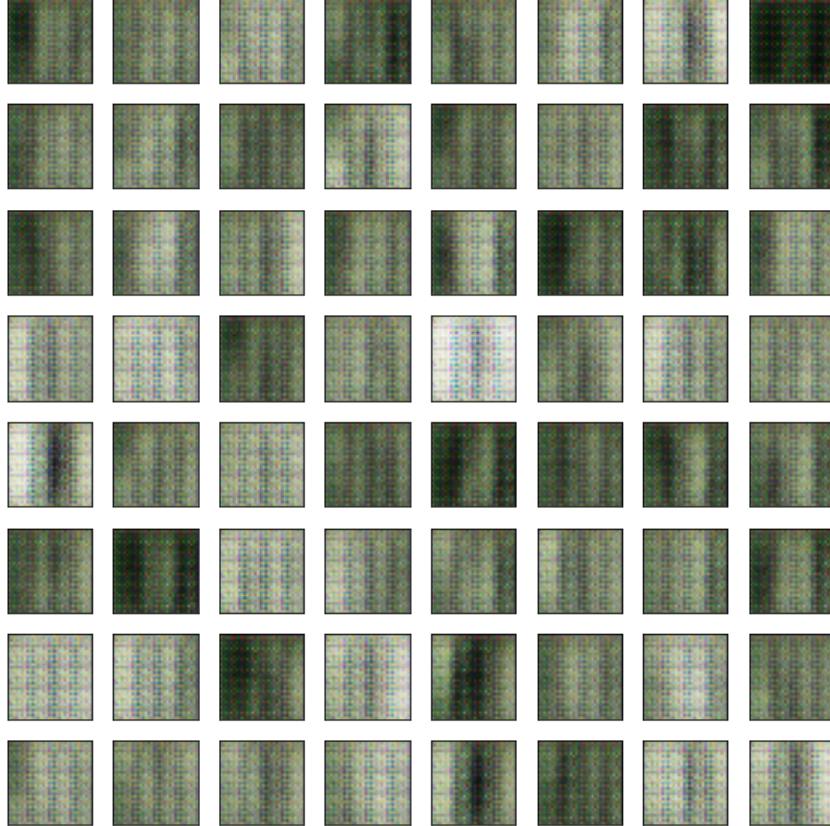
This is the 400/572 of epoch 1.

This is the 450/572 of epoch 1.

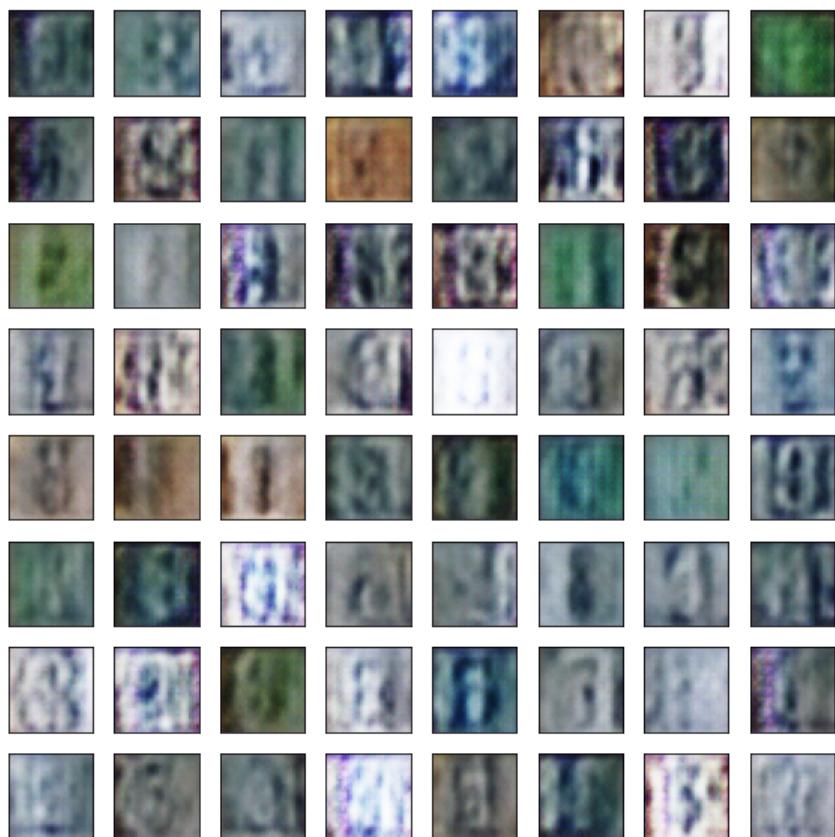
This is the 500/572 of epoch 1.

This is the 550/572 of epoch 1.

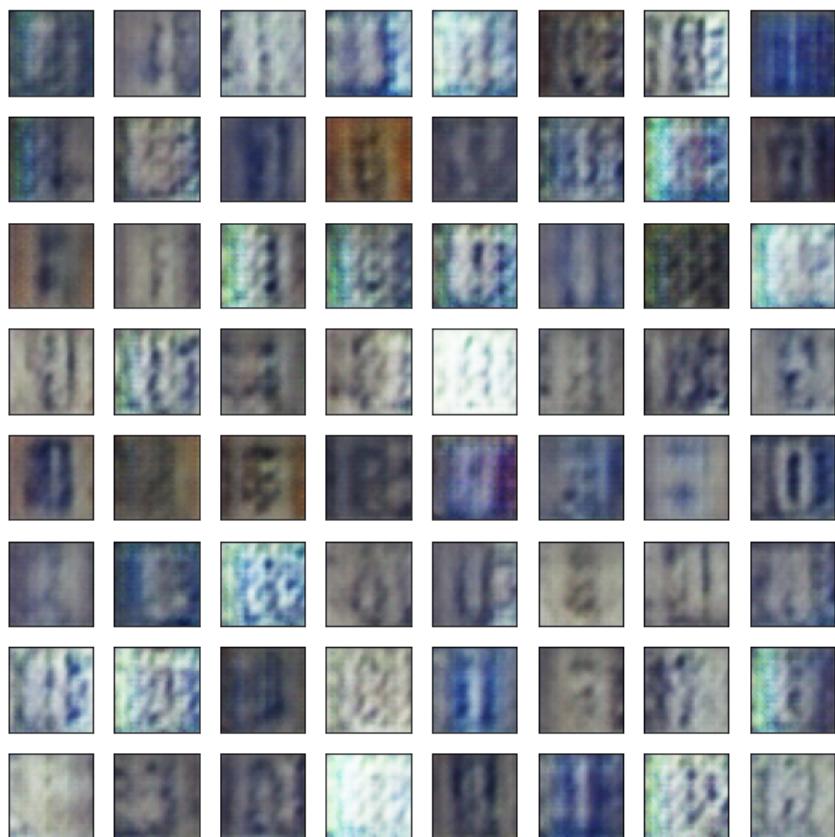
Epoch: 1/200 Discriminator Loss: 1.1818 Generator Loss: 2.2662



This is the 50/572 of epoch 2.  
This is the 100/572 of epoch 2.  
This is the 150/572 of epoch 2.  
This is the 200/572 of epoch 2.  
This is the 250/572 of epoch 2.  
This is the 300/572 of epoch 2.  
This is the 350/572 of epoch 2.  
This is the 400/572 of epoch 2.  
This is the 450/572 of epoch 2.  
This is the 500/572 of epoch 2.  
This is the 550/572 of epoch 2.  
Epoch: 2/200 Discriminator Loss: 1.0918 Generator Loss: 0.9048  
This is the 50/572 of epoch 3.  
This is the 100/572 of epoch 3.  
This is the 150/572 of epoch 3.  
This is the 200/572 of epoch 3.  
This is the 250/572 of epoch 3.  
This is the 300/572 of epoch 3.  
This is the 350/572 of epoch 3.  
This is the 400/572 of epoch 3.  
This is the 450/572 of epoch 3.  
This is the 500/572 of epoch 3.  
This is the 550/572 of epoch 3.  
Epoch: 3/200 Discriminator Loss: 1.0724 Generator Loss: 0.9790  
This is the 50/572 of epoch 4.  
This is the 100/572 of epoch 4.  
This is the 150/572 of epoch 4.  
This is the 200/572 of epoch 4.  
This is the 250/572 of epoch 4.  
This is the 300/572 of epoch 4.  
This is the 350/572 of epoch 4.  
This is the 400/572 of epoch 4.  
This is the 450/572 of epoch 4.  
This is the 500/572 of epoch 4.  
This is the 550/572 of epoch 4.  
Epoch: 4/200 Discriminator Loss: 1.6541 Generator Loss: 0.3343  
This is the 50/572 of epoch 5.  
This is the 100/572 of epoch 5.  
This is the 150/572 of epoch 5.  
This is the 200/572 of epoch 5.  
This is the 250/572 of epoch 5.  
This is the 300/572 of epoch 5.  
This is the 350/572 of epoch 5.  
This is the 400/572 of epoch 5.  
This is the 450/572 of epoch 5.  
This is the 500/572 of epoch 5.  
This is the 550/572 of epoch 5.  
Epoch: 5/200 Discriminator Loss: 1.4448 Generator Loss: 0.6570



This is the 50/572 of epoch 6.  
This is the 100/572 of epoch 6.  
This is the 150/572 of epoch 6.  
This is the 200/572 of epoch 6.  
This is the 250/572 of epoch 6.  
This is the 300/572 of epoch 6.  
This is the 350/572 of epoch 6.  
This is the 400/572 of epoch 6.  
This is the 450/572 of epoch 6.  
This is the 500/572 of epoch 6.  
This is the 550/572 of epoch 6.  
Epoch: 6/200 Discriminator Loss: 1.3173 Generator Loss: 1.4898  
This is the 50/572 of epoch 7.  
This is the 100/572 of epoch 7.  
This is the 150/572 of epoch 7.  
This is the 200/572 of epoch 7.  
This is the 250/572 of epoch 7.  
This is the 300/572 of epoch 7.  
This is the 350/572 of epoch 7.  
This is the 400/572 of epoch 7.  
This is the 450/572 of epoch 7.  
This is the 500/572 of epoch 7.  
This is the 550/572 of epoch 7.  
Epoch: 7/200 Discriminator Loss: 1.2405 Generator Loss: 1.3723  
This is the 50/572 of epoch 8.  
This is the 100/572 of epoch 8.  
This is the 150/572 of epoch 8.  
This is the 200/572 of epoch 8.  
This is the 250/572 of epoch 8.  
This is the 300/572 of epoch 8.  
This is the 350/572 of epoch 8.  
This is the 400/572 of epoch 8.  
This is the 450/572 of epoch 8.  
This is the 500/572 of epoch 8.  
This is the 550/572 of epoch 8.  
Epoch: 8/200 Discriminator Loss: 1.1848 Generator Loss: 0.8504  
This is the 50/572 of epoch 9.  
This is the 100/572 of epoch 9.  
This is the 150/572 of epoch 9.  
This is the 200/572 of epoch 9.  
This is the 250/572 of epoch 9.  
This is the 300/572 of epoch 9.  
This is the 350/572 of epoch 9.  
This is the 400/572 of epoch 9.  
This is the 450/572 of epoch 9.  
This is the 500/572 of epoch 9.  
This is the 550/572 of epoch 9.  
Epoch: 9/200 Discriminator Loss: 1.4371 Generator Loss: 1.2343  
This is the 50/572 of epoch 10.  
This is the 100/572 of epoch 10.  
This is the 150/572 of epoch 10.  
This is the 200/572 of epoch 10.  
This is the 250/572 of epoch 10.  
This is the 300/572 of epoch 10.  
This is the 350/572 of epoch 10.  
This is the 400/572 of epoch 10.  
This is the 450/572 of epoch 10.  
This is the 500/572 of epoch 10.  
This is the 550/572 of epoch 10.  
Epoch: 10/200 Discriminator Loss: 1.5515 Generator Loss: 0.5687



This is the 50/572 of epoch 46.  
This is the 100/572 of epoch 46.  
This is the 150/572 of epoch 46.  
This is the 200/572 of epoch 46.  
This is the 250/572 of epoch 46.  
This is the 300/572 of epoch 46.  
This is the 350/572 of epoch 46.  
This is the 400/572 of epoch 46.  
This is the 450/572 of epoch 46.  
This is the 500/572 of epoch 46.  
This is the 550/572 of epoch 46.  
Epoch: 46/200 Discriminator Loss: 1.5772 Generator Loss: 2.1724  
This is the 50/572 of epoch 47.  
This is the 100/572 of epoch 47.  
This is the 150/572 of epoch 47.  
This is the 200/572 of epoch 47.  
This is the 250/572 of epoch 47.  
This is the 300/572 of epoch 47.  
This is the 350/572 of epoch 47.  
This is the 400/572 of epoch 47.  
This is the 450/572 of epoch 47.  
This is the 500/572 of epoch 47.  
This is the 550/572 of epoch 47.  
Epoch: 47/200 Discriminator Loss: 1.3849 Generator Loss: 2.0000  
This is the 50/572 of epoch 48.  
This is the 100/572 of epoch 48.  
This is the 150/572 of epoch 48.  
This is the 200/572 of epoch 48.  
This is the 250/572 of epoch 48.  
This is the 300/572 of epoch 48.  
This is the 350/572 of epoch 48.  
This is the 400/572 of epoch 48.  
This is the 450/572 of epoch 48.  
This is the 500/572 of epoch 48.  
This is the 550/572 of epoch 48.  
Epoch: 48/200 Discriminator Loss: 1.4218 Generator Loss: 2.3933  
This is the 50/572 of epoch 49.  
This is the 100/572 of epoch 49.  
This is the 150/572 of epoch 49.  
This is the 200/572 of epoch 49.  
This is the 250/572 of epoch 49.  
This is the 300/572 of epoch 49.  
This is the 350/572 of epoch 49.  
This is the 400/572 of epoch 49.  
This is the 450/572 of epoch 49.  
This is the 500/572 of epoch 49.  
This is the 550/572 of epoch 49.  
Epoch: 49/200 Discriminator Loss: 1.4778 Generator Loss: 1.6668  
This is the 50/572 of epoch 50.  
This is the 100/572 of epoch 50.  
This is the 150/572 of epoch 50.  
This is the 200/572 of epoch 50.  
This is the 250/572 of epoch 50.  
This is the 300/572 of epoch 50.  
This is the 350/572 of epoch 50.  
This is the 400/572 of epoch 50.  
This is the 450/572 of epoch 50.  
This is the 500/572 of epoch 50.  
This is the 550/572 of epoch 50.  
Epoch: 50/200 Discriminator Loss: 1.4739 Generator Loss: 2.1755



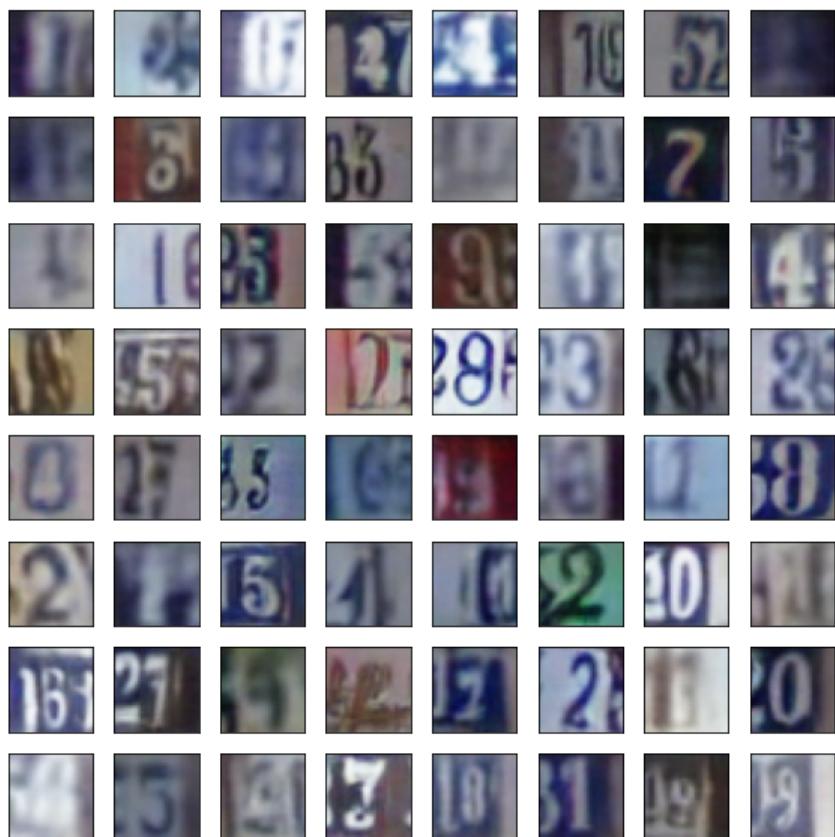
This is the 50/572 of epoch 96.  
This is the 100/572 of epoch 96.  
This is the 150/572 of epoch 96.  
This is the 200/572 of epoch 96.  
This is the 250/572 of epoch 96.  
This is the 300/572 of epoch 96.  
This is the 350/572 of epoch 96.  
This is the 400/572 of epoch 96.  
This is the 450/572 of epoch 96.  
This is the 500/572 of epoch 96.  
This is the 550/572 of epoch 96.  
Epoch: 96/200 Discriminator Loss: 1.5604 Generator Loss: 2.1650  
This is the 50/572 of epoch 97.  
This is the 100/572 of epoch 97.  
This is the 150/572 of epoch 97.  
This is the 200/572 of epoch 97.  
This is the 250/572 of epoch 97.  
This is the 300/572 of epoch 97.  
This is the 350/572 of epoch 97.  
This is the 400/572 of epoch 97.  
This is the 450/572 of epoch 97.  
This is the 500/572 of epoch 97.  
This is the 550/572 of epoch 97.  
Epoch: 97/200 Discriminator Loss: 1.9035 Generator Loss: 2.6662  
This is the 50/572 of epoch 98.  
This is the 100/572 of epoch 98.  
This is the 150/572 of epoch 98.  
This is the 200/572 of epoch 98.  
This is the 250/572 of epoch 98.  
This is the 300/572 of epoch 98.  
This is the 350/572 of epoch 98.  
This is the 400/572 of epoch 98.  
This is the 450/572 of epoch 98.  
This is the 500/572 of epoch 98.  
This is the 550/572 of epoch 98.  
Epoch: 98/200 Discriminator Loss: 1.4342 Generator Loss: 1.9297  
This is the 50/572 of epoch 99.  
This is the 100/572 of epoch 99.  
This is the 150/572 of epoch 99.  
This is the 200/572 of epoch 99.  
This is the 250/572 of epoch 99.  
This is the 300/572 of epoch 99.  
This is the 350/572 of epoch 99.  
This is the 400/572 of epoch 99.  
This is the 450/572 of epoch 99.  
This is the 500/572 of epoch 99.  
This is the 550/572 of epoch 99.  
Epoch: 99/200 Discriminator Loss: 1.1888 Generator Loss: 1.0061  
This is the 50/572 of epoch 100.  
This is the 100/572 of epoch 100.  
This is the 150/572 of epoch 100.  
This is the 200/572 of epoch 100.  
This is the 250/572 of epoch 100.  
This is the 300/572 of epoch 100.  
This is the 350/572 of epoch 100.  
This is the 400/572 of epoch 100.  
This is the 450/572 of epoch 100.  
This is the 500/572 of epoch 100.  
This is the 550/572 of epoch 100.  
Epoch: 100/200 Discriminator Loss: 1.6090 Generator Loss: 2.0289



This is the 50/572 of epoch 146.  
This is the 100/572 of epoch 146.  
This is the 150/572 of epoch 146.  
This is the 200/572 of epoch 146.  
This is the 250/572 of epoch 146.  
This is the 300/572 of epoch 146.  
This is the 350/572 of epoch 146.  
This is the 400/572 of epoch 146.  
This is the 450/572 of epoch 146.  
This is the 500/572 of epoch 146.  
This is the 550/572 of epoch 146.  
Epoch: 146/200 Discriminator Loss: 2.4261 Generator Loss: 2.7198  
This is the 50/572 of epoch 147.  
This is the 100/572 of epoch 147.  
This is the 150/572 of epoch 147.  
This is the 200/572 of epoch 147.  
This is the 250/572 of epoch 147.  
This is the 300/572 of epoch 147.  
This is the 350/572 of epoch 147.  
This is the 400/572 of epoch 147.  
This is the 450/572 of epoch 147.  
This is the 500/572 of epoch 147.  
This is the 550/572 of epoch 147.  
Epoch: 147/200 Discriminator Loss: 1.6256 Generator Loss: 2.2680  
This is the 50/572 of epoch 148.  
This is the 100/572 of epoch 148.  
This is the 150/572 of epoch 148.  
This is the 200/572 of epoch 148.  
This is the 250/572 of epoch 148.  
This is the 300/572 of epoch 148.  
This is the 350/572 of epoch 148.  
This is the 400/572 of epoch 148.  
This is the 450/572 of epoch 148.  
This is the 500/572 of epoch 148.  
This is the 550/572 of epoch 148.  
Epoch: 148/200 Discriminator Loss: 1.4876 Generator Loss: 1.7390  
This is the 50/572 of epoch 149.  
This is the 100/572 of epoch 149.  
This is the 150/572 of epoch 149.  
This is the 200/572 of epoch 149.  
This is the 250/572 of epoch 149.  
This is the 300/572 of epoch 149.  
This is the 350/572 of epoch 149.  
This is the 400/572 of epoch 149.  
This is the 450/572 of epoch 149.  
This is the 500/572 of epoch 149.  
This is the 550/572 of epoch 149.  
Epoch: 149/200 Discriminator Loss: 1.6378 Generator Loss: 1.9026  
This is the 50/572 of epoch 150.  
This is the 100/572 of epoch 150.  
This is the 150/572 of epoch 150.  
This is the 200/572 of epoch 150.  
This is the 250/572 of epoch 150.  
This is the 300/572 of epoch 150.  
This is the 350/572 of epoch 150.  
This is the 400/572 of epoch 150.  
This is the 450/572 of epoch 150.  
This is the 500/572 of epoch 150.  
This is the 550/572 of epoch 150.  
Epoch: 150/200 Discriminator Loss: 1.5081 Generator Loss: 1.7136



This is the 50/572 of epoch 196.  
This is the 100/572 of epoch 196.  
This is the 150/572 of epoch 196.  
This is the 200/572 of epoch 196.  
This is the 250/572 of epoch 196.  
This is the 300/572 of epoch 196.  
This is the 350/572 of epoch 196.  
This is the 400/572 of epoch 196.  
This is the 450/572 of epoch 196.  
This is the 500/572 of epoch 196.  
This is the 550/572 of epoch 196.  
Epoch: 196/200 Discriminator Loss: 1.7758 Generator Loss: 2.0924  
This is the 50/572 of epoch 197.  
This is the 100/572 of epoch 197.  
This is the 150/572 of epoch 197.  
This is the 200/572 of epoch 197.  
This is the 250/572 of epoch 197.  
This is the 300/572 of epoch 197.  
This is the 350/572 of epoch 197.  
This is the 400/572 of epoch 197.  
This is the 450/572 of epoch 197.  
This is the 500/572 of epoch 197.  
This is the 550/572 of epoch 197.  
Epoch: 197/200 Discriminator Loss: 1.8074 Generator Loss: 2.2938  
This is the 50/572 of epoch 198.  
This is the 100/572 of epoch 198.  
This is the 150/572 of epoch 198.  
This is the 200/572 of epoch 198.  
This is the 250/572 of epoch 198.  
This is the 300/572 of epoch 198.  
This is the 350/572 of epoch 198.  
This is the 400/572 of epoch 198.  
This is the 450/572 of epoch 198.  
This is the 500/572 of epoch 198.  
This is the 550/572 of epoch 198.  
Epoch: 198/200 Discriminator Loss: 1.8307 Generator Loss: 1.8116  
This is the 50/572 of epoch 199.  
This is the 100/572 of epoch 199.  
This is the 150/572 of epoch 199.  
This is the 200/572 of epoch 199.  
This is the 250/572 of epoch 199.  
This is the 300/572 of epoch 199.  
This is the 350/572 of epoch 199.  
This is the 400/572 of epoch 199.  
This is the 450/572 of epoch 199.  
This is the 500/572 of epoch 199.  
This is the 550/572 of epoch 199.  
Epoch: 199/200 Discriminator Loss: 2.2312 Generator Loss: 2.7710  
This is the 50/572 of epoch 200.  
This is the 100/572 of epoch 200.  
This is the 150/572 of epoch 200.  
This is the 200/572 of epoch 200.  
This is the 250/572 of epoch 200.  
This is the 300/572 of epoch 200.  
This is the 350/572 of epoch 200.  
This is the 400/572 of epoch 200.  
This is the 450/572 of epoch 200.  
This is the 500/572 of epoch 200.  
This is the 550/572 of epoch 200.  
Epoch: 200/200 Discriminator Loss: 1.4912 Generator Loss: 1.7491



12572.459484577179

Out[17]: <tensorflow.python.keras.engine.sequential.Sequential at 0x7ff5f0185240>

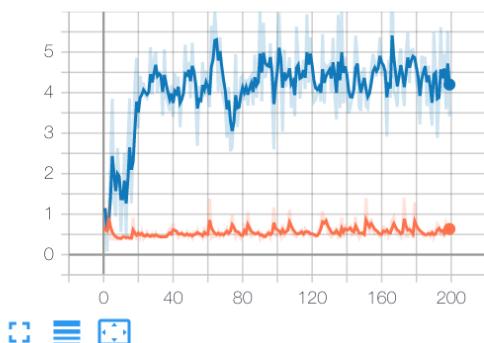
In [18]: %tensorboard --logdir=~/logs\_SVHN

---

epoch\_training\_loss

---

epoch\_training\_loss



---

epoch\_validation\_loss

---

epoch\_validation\_loss

