

```
In [ ]: import matplotlib.pyplot as plt
import numpy as np
import os
from tensorflow.keras import layers
import time
import tensorflow as tf
from IPython import display
import datetime
import pickle as pkl
from scipy.io import loadmat
from tensorflow.keras.utils import get_file
```

```
In [ ]: %load_ext tensorboard
!rm -rf ./logs/wgan_gradient_tape/
```

The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

```
In [ ]: data_file_train = get_file('train_32x32.mat', origin='http://ufldl.stanford.edu/housenumbers/train_32x32.mat')
data_file_valid = get_file('test_32x32.mat', origin='http://ufldl.stanford.edu/housenumbers/test_32x32.mat')

train_dir = os.path.join(os.path.dirname(data_file_train), 'train_32x32.mat')
valid_dir = os.path.join(os.path.dirname(data_file_valid), 'test_32x32.mat')

train_data = loadmat(train_dir)
valid_data = loadmat(valid_dir)

X_train, y_train = train_data['X'], train_data['y']
X_valid, y_valid = valid_data['X'], valid_data['y']

print(X_train.shape) # shape is in (image_height, image_width, channels, records)
```

```
Downloading data from http://ufldl.stanford.edu/housenumbers/train_32x32.mat
182042624/182040794 [=====] - 6s 0us/step
Downloading data from http://ufldl.stanford.edu/housenumbers/test_32x32.mat
64282624/64275384 [=====] - 3s 0us/step
(32, 32, 3, 73257)
```

```
In [ ]: X_train = np.rollaxis(X_train, 3)
X_valid = np.rollaxis(X_valid, 3)

print(X_train.shape) # convert shape to (records, image_height, image_width, channels)
print(X_valid.shape)

(73257, 32, 32, 3)
(26032, 32, 32, 3)
```

```
In [ ]: def preprocess(x):
    return (x-127.5)/127.5 # standardize to [-1, 1], so that tanh function is applicable

def decode(x):
    return np.uint8(127.5*x+127.5) # make sure to use uint8 type otherwise the image won't display properly
```

```
In [ ]: X_train_real = preprocess(X_train)
X_valid_real = preprocess(X_valid)
```

```
In [ ]: # glimpse of the real images
sample_images = X_train[np.random.choice(len(X_train_real), size=64, replace=False
)]

plt.figure(figsize=(8, 8))
for i in range(64):
    plt.subplot(8, 8, i+1)
    plt.imshow(sample_images[i])
    plt.xticks([])
    plt.yticks([])
plt.tight_layout()
plt.show()
```



```
In [ ]: batch_size = 64
n_critic = 1
#n_critic = 1
num_batches = len(X_train_real)//batch_size
```

```
In [ ]: def build_generator(input_size, leaky_alpha=0.2):
        model = tf.keras.Sequential([
            layers.Dense(4*4*512, input_shape=(input_size,)),
            layers.Reshape(target_shape=(4, 4, 512)), # 4,
            4,512
            layers.BatchNormalization(),
            layers.LeakyReLU(alpha=leaky_alpha),
            layers.Dropout(0.2),
            layers.Conv2DTranspose(256, kernel_size=5, strides=2, padding='same', use_bi
as=False), # 8,8,256
            layers.BatchNormalization(),
            layers.LeakyReLU(alpha=leaky_alpha),
            layers.Conv2DTranspose(128, kernel_size=5, strides=2, padding='same', use_bi
as=False), # 16,16,128
            layers.BatchNormalization(),
            layers.LeakyReLU(alpha=leaky_alpha),
            layers.Conv2DTranspose(3, kernel_size=5, strides=2, padding='same', use_bias
=False), # 32,32,3
            layers.Activation('tanh')
        ])

        return model
```

```
In [ ]: generator = build_generator(100, 0.2)
        generator.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 8192)	827392
reshape (Reshape)	(None, 4, 4, 512)	0
batch_normalization (BatchNo	(None, 4, 4, 512)	2048
leaky_re_lu (LeakyReLU)	(None, 4, 4, 512)	0
dropout (Dropout)	(None, 4, 4, 512)	0
conv2d_transpose (Conv2DTran	(None, 8, 8, 256)	3276800
batch_normalization_1 (Batch	(None, 8, 8, 256)	1024
leaky_re_lu_1 (LeakyReLU)	(None, 8, 8, 256)	0
conv2d_transpose_1 (Conv2DTr	(None, 16, 16, 128)	819200
batch_normalization_2 (Batch	(None, 16, 16, 128)	512
leaky_re_lu_2 (LeakyReLU)	(None, 16, 16, 128)	0
conv2d_transpose_2 (Conv2DTr	(None, 32, 32, 3)	9600
activation (Activation)	(None, 32, 32, 3)	0
Total params: 4,936,576		
Trainable params: 4,934,784		
Non-trainable params: 1,792		

```
In [ ]: def build_discriminator(leaky_alpha=0.2):
        model = tf.keras.Sequential([
            layers.Conv2D(64, (5,5), strides=2, padding='same', input_shape=(32,32,3
        )), # 16,16,64
            layers.LeakyReLU(leaky_alpha),
            layers.Dropout(0.2),
            layers.Conv2D(128, (5,5), strides=2, padding='same'), # 8,8,128
            layers.BatchNormalization(),
            layers.LeakyReLU(leaky_alpha),
            layers.Dropout(0.2),
            layers.Conv2D(256, (5,5), strides=2, padding='same'), # 4,4,256
            layers.BatchNormalization(),
            layers.LeakyReLU(leaky_alpha),
            layers.Flatten(),
            layers.Dense(1)
        ])
        return model
```

```
In [ ]: discriminator = build_discriminator(0.2)
discriminator.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 16, 16, 64)	4864
leaky_re_lu_3 (LeakyReLU)	(None, 16, 16, 64)	0
dropout_1 (Dropout)	(None, 16, 16, 64)	0
conv2d_1 (Conv2D)	(None, 8, 8, 128)	204928
batch_normalization_3 (Batch Normalization)	(None, 8, 8, 128)	512
leaky_re_lu_4 (LeakyReLU)	(None, 8, 8, 128)	0
dropout_2 (Dropout)	(None, 8, 8, 128)	0
conv2d_2 (Conv2D)	(None, 4, 4, 256)	819456
batch_normalization_4 (Batch Normalization)	(None, 4, 4, 256)	1024
leaky_re_lu_5 (LeakyReLU)	(None, 4, 4, 256)	0
flatten (Flatten)	(None, 4096)	0
dense_1 (Dense)	(None, 1)	4097
=====		
Total params: 1,034,881		
Trainable params: 1,034,113		
Non-trainable params: 768		

```
In [ ]: generator_optimizer = tf.keras.optimizers.RMSprop(learning_rate=5e-5)
discriminator_optimizer = tf.keras.optimizers.RMSprop(learning_rate=5e-5)
```

```
In [ ]: def show_images(generated_images):
    n_images = len(generated_images)
    cols = 8
    rows = n_images//cols

    plt.figure(figsize=(8, 8))
    for i in range(n_images):
        img = decode(generated_images[i])
        ax = plt.subplot(rows, cols, i+1)
        plt.imshow(img)
        plt.xticks([])
        plt.yticks([])
    plt.tight_layout()
    plt.show()
```

```
In [ ]: EPOCHS = 200
np.random.seed(2020)
seed = np.random.normal(loc=0, scale=1, size=(64, 100))

current_time = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
gen_log_dir = 'logs/wgan_gradient_tape/' + current_time + '/gen'
disc_log_dir = 'logs/wgan_gradient_tape/' + current_time + '/disc'
gen_summary_writer = tf.summary.create_file_writer(gen_log_dir)
disc_summary_writer = tf.summary.create_file_writer(disc_log_dir)
```

```
In [ ]: @tf.function
def generator_gradient():
    # step 9
    noise = tf.random.normal([batch_size, 100])
    with tf.GradientTape() as gen_tape:
        # step 10
        generated_images = generator(noise, training=True)
        gen_loss_val = -tf.reduce_mean(discriminator(generated_images, training=True))
    # step 10
    gradients_of_generator = gen_tape.gradient(gen_loss_val, generator.trainable_variables)
    # step 11
    generator_optimizer.apply_gradients(zip(gradients_of_generator, generator.trainable_variables))
    return gen_loss_val
```

```

In [ ]: @tf.function
def discriminator_gradient():
    # step 3
    noise = tf.random.normal([batch_size, 100])
    # step 4
    images = X_train_real[np.random.choice(len(X_train_real), size=batch_size, replace=False)]

    with tf.GradientTape() as disc_tape:
        # step 5
        generated_images = generator(noise, training=True)
        dis_loss_val = -tf.reduce_mean(discriminator(images, training=True)) + tf.reduce_mean(discriminator(generated_images, training=True))
        # step 5
        gradients_of_discriminator = disc_tape.gradient(dis_loss_val, discriminator.trainable_variables)
        # step 6
        discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator, discriminator.trainable_variables))

    # step 7, clip
    for p in discriminator.trainable_variables:
        p.assign(tf.clip_by_value(p, -0.01, 0.01))

    return dis_loss_val

```

```

In [ ]: # the number of step corresponds to Algorithm 1 in "Wasserstein GAN"
def train(epochs):

    start = time.time()
    # step 1
    for e in range(epochs):
        start2 = time.time()
        for i in range(num_batches):
            # step 2
            for _ in range(n_critic):
                dis_loss_val = discriminator_gradient()
            # step 8
            gen_loss_val = generator_gradient()
            #if (i+1) % 50 == 0:
                #print('This is the {}/{} of epoch {}'.format(i+1, num_batches, e+1))

            with gen_summary_writer.as_default():
                tf.summary.scalar('loss', gen_loss_val, step=e)
            with disc_summary_writer.as_default():
                tf.summary.scalar('loss', dis_loss_val, step=e)


            print("Epoch: {}/{} Discriminator Loss: {:.4f} Generator Loss: {:.4f}".format(e+1, epochs, dis_loss_val, gen_loss_val))
            if e==0 or e==4 or e==9 or e==49 or e==99 or e==149 or e==199:
                imgs = generator.predict_on_batch(seed)
                show_images(imgs)
            print('Time for epoch {} is {} sec'.format(e + 1, time.time()-start2))

        # step 12
        end = time.time()
        print(end-start)

```

```
In [ ]: train(200)
```

Epoch: 1/200 Discriminator Loss: -0.0149 Generator Loss: -0.0062



Time for epoch 1 is 61.050485134124756 sec

Epoch: 2/200 Discriminator Loss: -0.0136 Generator Loss: -0.0073

Time for epoch 2 is 56.70659685134888 sec


Epoch: 3/200 Discriminator Loss: -0.0137 Generator Loss: -0.0053

Time for epoch 3 is 57.16050457954407 sec

Epoch: 4/200 Discriminator Loss: -0.0162 Generator Loss: -0.0072

Time for epoch 4 is 57.761735677719116 sec

Epoch: 5/200 Discriminator Loss: -0.0131 Generator Loss: -0.0059



Time for epoch 5 is 59.584129095077515 sec

Epoch: 6/200 Discriminator Loss: -0.0137 Generator Loss: -0.0020

Time for epoch 6 is 57.57205772399902 sec

Epoch: 7/200 Discriminator Loss: -0.0112 Generator Loss: -0.0067

Time for epoch 7 is 57.804524183273315 sec

Epoch: 8/200 Discriminator Loss: -0.0135 Generator Loss: -0.0077

Time for epoch 8 is 57.805200815200806 sec

Epoch: 9/200 Discriminator Loss: -0.0126 Generator Loss: -0.0069

Time for epoch 9 is 57.811567068099976 sec

Epoch: 10/200 Discriminator Loss: -0.0120 Generator Loss: -0.0083



Epoch: 42/200 Discriminator Loss: -0.0096 Generator Loss: -0.0099
Time for epoch 42 is 57.61907720565796 sec
Epoch: 43/200 Discriminator Loss: -0.0096 Generator Loss: -0.0091
Time for epoch 43 is 57.65308117866516 sec
Epoch: 44/200 Discriminator Loss: -0.0115 Generator Loss: -0.0115
Time for epoch 44 is 57.54136562347412 sec
Epoch: 45/200 Discriminator Loss: -0.0093 Generator Loss: -0.0074
Time for epoch 45 is 57.62323474884033 sec
Epoch: 46/200 Discriminator Loss: -0.0096 Generator Loss: -0.0088
Time for epoch 46 is 57.357157945632935 sec
Epoch: 47/200 Discriminator Loss: -0.0099 Generator Loss: -0.0100
Time for epoch 47 is 57.439390659332275 sec
Epoch: 48/200 Discriminator Loss: -0.0100 Generator Loss: -0.0094
Time for epoch 48 is 57.452401876449585 sec
Epoch: 49/200 Discriminator Loss: -0.0104 Generator Loss: -0.0098
Time for epoch 49 is 57.338640213012695 sec
Epoch: 50/200 Discriminator Loss: -0.0083 Generator Loss: -0.0078



Time for epoch 50 is 59.48170065879822 sec
Epoch: 51/200 Discriminator Loss: -0.0074 Generator Loss: -0.0101
Time for epoch 51 is 57.458508014678955 sec
Epoch: 52/200 Discriminator Loss: -0.0102 Generator Loss: -0.0075
Time for epoch 52 is 57.83629894256592 sec
Epoch: 54/200 Discriminator Loss: -0.0105 Generator Loss: -0.0114
Time for epoch 54 is 57.316943407058716 sec

In []: %tensorboard --logdir logs/wgan_gradient_tape

loss

loss
tag: loss

